

TITLE

July 19, 2026

CONTENTS

1. Elementary Foundations (without measure-theoretic language)	2
2. Measure-Theoretic Foundations	5
3. Independence (Measure-Theoretic)	6
4. Moments and Joint Distributions	7
5. Inequalities	8
6. Modes of Convergence	10
7. Limit Theorems	11
8. Integration Theorems (used in Probability)	12
9. Conditional Expectation	13
10. Characteristic Functions	14
11. Common Probability Distributions	14
11.1. Discrete Distributions	14
11.2. Continuous Distributions	15
11.3. Multivariate Distributions	17
12. Central Limit Theorem	17
13. Martingales	18
14. Stochastic Processes	20
15. Entropy and Information Theory	20
16. Conditional Probability and Bayes' Theorem	22
17. Moment Generating Functions	23
18. Coupling and Total Variation Distance	23
19. Ergodic Theory (Probability-Theoretic View)	24

1. ELEMENTARY FOUNDATIONS (WITHOUT MEASURE-THEORETIC LANGUAGE)

Definition 1.0.1 (Probability space). A *probability space* is a triple $(\Omega, \mathcal{F}, P)$ consisting of:

1. A nonempty set Ω , called the *sample space*, whose elements are the possible outcomes of a random experiment.
2. A collection $\mathcal{F}$ of subsets of Ω , called the *events*, such that:
 - (a) $\Omega \in \mathcal{F}$,
 - (b) if $E \in \mathcal{F}$ then its complement $E^c = \Omega \setminus E$ is in $\mathcal{F}$,
 - (c) if $(E_n)_{n=1}^\infty$ is a sequence with each $E_n \in \mathcal{F}$, then $\bigcup_{n=1}^\infty E_n \in \mathcal{F}$.
3. A map $P : \mathcal{F} \rightarrow [0, 1]$, called the *probability measure*, satisfying:
 - (a) $P(\Omega) = 1$,
 - (b) for every sequence $(E_n)_{n=1}^\infty$ of pairwise disjoint events (i.e., $E_i \cap E_j = \emptyset$ for $i \neq j$),

$$P\left(\bigcup_{n=1}^\infty E_n\right) = \sum_{n=1}^\infty P(E_n).$$

The value $P(E)$ is called the *probability of the event E* .

Definition 1.0.2 (Random variable). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1).

A function $X : \Omega \rightarrow S$ into a set S equipped with a collection $\mathcal{S}$ of subsets of S is called a *$(S, \mathcal{S})$ -valued random variable* if for every $A \in \mathcal{S}$, the set

$$X^{-1}(A) = \{\omega \in \Omega : X(\omega) \in A\}$$

is an event in $\mathcal{F}$.

When S is a countable or finite set and $\mathcal{S}$ is the power set of S (i.e., every subset of S belongs to $\mathcal{S}$), then X is a random variable precisely when $X^{-1}(\{s\}) \in \mathcal{F}$ for every $s \in S$. This is the case that arises in introductory probability theory, where one studies discrete random variables taking values in sets such as $\{0, 1\}$ (Bernoulli trials), $\{0, 1, 2, \dots\}$ (counting processes), or any other countable or finite sample space.

Special cases:

1. A function $X : \Omega \rightarrow \mathbb{R}$ is called a *real-valued random variable* if for every pair of real numbers $a < b$, the set

$$\{\omega \in \Omega : X(\omega) \in (a, b)\} = \{\omega \in \Omega : a < X(\omega) < b\}$$

is an event in $\mathcal{F}$.

2. A function $X : \Omega \rightarrow \mathbb{C}$ is called a *complex-valued random variable* if for every open disk $D \subseteq \mathbb{C}$, the set

$$\{\omega \in \Omega : X(\omega) \in D\}$$

is an event in $\mathcal{F}$.

3. A function $X : \Omega \rightarrow \overline{\mathbb{R}} = \mathbb{R} \cup \{-\infty, \infty\}$ is called a *random variable taking values in the extended real line* if for every $a \in \mathbb{R}$, the set

$$\{\omega \in \Omega : X(\omega) > a\}$$

is an event in $\mathcal{F}$.

A function X as above ($(S, \mathcal{S})$ -valued, real-valued, complex-valued, or extended real-valued) is called simply a **random variable** when the codomain and collection of subsets are clear from context.

Definition 1.0.3 (Distribution (law) of a random variable on a probability space). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1), and let $X : \Omega \rightarrow S$ be a random variable taking values in a set S .

The **distribution** (or **law**) of X , denoted μ_X , describes how the probability mass of P is distributed over S through X : for any event $A \subseteq S$ whose occurrence makes sense in context,

$$\mu_X(A) = P(X \in A).$$

In the special case where $S = \mathbb{R}$, the distribution of X is uniquely determined by its cumulative distribution function $F_X : \mathbb{R} \rightarrow [0, 1]$ given by

$$F_X(t) = P(X \leq t).$$

Definition 1.0.4 (Expectation of a random variable on a probability space). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1).

1. If $X : \Omega \rightarrow [0, \infty)$ is a nonnegative random variable, its **expectation** (or **expected value**) $\mathbb{E}[X]$ is the average value of X over all outcomes, weighted by their probabilities.
2. If $X : \Omega \rightarrow \mathbb{R}$ is a real-valued random variable, write $X = X^+ - X^-$ where $X^+ = \max(X, 0)$ and $X^- = \max(-X, 0)$. The expectation of X is defined by

$$\mathbb{E}[X] := \mathbb{E}[X^+] - \mathbb{E}[X^-],$$

provided that at least one of $\mathbb{E}[X^+]$, $\mathbb{E}[X^-]$ is finite. If both are finite (i.e., X is integrable), then $\mathbb{E}[X]$ is a well-defined real number.

3. A random variable X is called **integrable** if $\mathbb{E}[|X|] < \infty$. The space of all integrable real-valued random variables on $(\Omega, \mathcal{F}, P)$ is denoted by $L^1(\Omega, \mathcal{F}, P)$.

If $X : \Omega \rightarrow S$ is an $(S, \mathcal{S})$ -valued random variable with distribution μ_X on $(S, \mathcal{S})$, and if $g : S \rightarrow \mathbb{R}$ is a function such that $g(X)$ is integrable, then the **change of variables formula** holds:

$$\mathbb{E}[g(X)] = \int_{\Omega} g(X(\omega)) dP(\omega) = \int_S g(x) d\mu_X(x).$$

Definition 1.0.5 (Cumulative distribution function). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1). Let $X : \Omega \rightarrow \mathbb{R}$ be a function such that for every $t \in \mathbb{R}$, the set $\{\omega \in \Omega : X(\omega) \leq t\}$ belongs to $\mathcal{F}$. The **cumulative distribution function of X** , denoted by F_X , is the function $F_X : \mathbb{R} \rightarrow [0, 1]$ defined by

$$F_X(t) = P(\{\omega \in \Omega : X(\omega) \leq t\}).$$

Definition 1.0.6 (Product probability space). Let $(\Omega_1, \mathcal{F}_1, P_1)$ and $(\Omega_2, \mathcal{F}_2, P_2)$ be two probability measure spaces (Definition 2.0.1).

The **product probability space** is the triple $(\Omega_1 \times \Omega_2, \mathcal{F}_1 \otimes \mathcal{F}_2, P_1 \times P_2)$ where:

1. $\mathcal{F}_1 \otimes \mathcal{F}_2$ is the product σ -algebra on $\Omega_1 \times \Omega_2$, and
2. $P_1 \times P_2$ is the product measure, which satisfies $(P_1 \times P_2)(A_1 \times A_2) = P_1(A_1)P_2(A_2)$ for all $A_1 \in \mathcal{F}_1$ and $A_2 \in \mathcal{F}_2$.

Since $P_1(\Omega_1) = P_2(\Omega_2) = 1$, the product measure satisfies $(P_1 \times P_2)(\Omega_1 \times \Omega_2) = 1$, so the product is again a probability space. This construction extends inductively to finite products and, via the Kolmogorov extension theorem, to countably infinite products.

Definition 1.0.7 (Independence of events). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1).

1. A finite collection of events $A_1, A_2, \dots, A_n \in \mathcal{F}$ is called **independent** if for every nonempty subset $I \subseteq \{1, 2, \dots, n\}$,

$$P\left(\bigcap_{i \in I} A_i\right) = \prod_{i \in I} P(A_i).$$

2. An arbitrary (possibly infinite) collection of events $\mathcal{C} \subseteq \mathcal{F}$ is called **independent** if every finite subcollection of $\mathcal{C}$ is independent in the sense of item 1. That is, for any distinct $A_1, A_2, \dots, A_n \in \mathcal{C}$,

$$P\left(\bigcap_{i=1}^n A_i\right) = \prod_{i=1}^n P(A_i).$$

A sequence $(A_n)_{n=1}^\infty$ of events is called an **independent sequence** if the set $\{A_1, A_2, \dots\}$ is independent.

Definition 1.0.8 (Independence of random variables on a probability space). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1).

A finite collection of real-valued random variables $X_1, X_2, \dots, X_n$ on $(\Omega, \mathcal{F}, P)$ is called **independent** if for every choice of intervals $I_1, I_2, \dots, I_n \subseteq \mathbb{R}$,

$$P(\{\omega : X_1(\omega) \in I_1\} \cap \{\omega : X_2(\omega) \in I_2\} \cap \dots \cap \{\omega : X_n(\omega) \in I_n\}) = P(X_1 \in I_1)P(X_2 \in I_2) \cdots P(X_n \in I_n)$$

An arbitrary (possibly infinite) collection of random variables $\{X_j\}_{j \in J}$ is called **independent** if every finite subcollection is independent.

A sequence $(X_n)_{n=1}^\infty$ of random variables is called an **independent sequence** if every finite subcollection is independent.

Proposition 1.0.9 (Jensen's inequality). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1). Let $\varphi : I \rightarrow \mathbb{R}$ be a convex function defined on an interval $I \subseteq \mathbb{R}$. If $X : \Omega \rightarrow I$ is a random variable such that $\mathbb{E}[|X|] < \infty$ and $\mathbb{E}[|\varphi(X)|] < \infty$, then

$$\varphi(\mathbb{E}[X]) \leq \mathbb{E}[\varphi(X)].$$

2. MEASURE-THEORETIC FOUNDATIONS

Definition 2.0.1 (Probability measure space). A measure space $(\Omega, \mathcal{F}, P)$ is called a *probability measure space*, and P is called a *probability measure* if $P(\Omega) = 1$.

In this case, Ω is called the *space of outcomes*, elements of $\mathcal{F}$ are called *events* and P is called the *probability measure*.

Definition 2.0.2 (Random variable on a probability measure space). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1).

A function $X : \Omega \rightarrow S$ into a measurable space $(S, \mathcal{S})$ is called a *$(S, \mathcal{S})$ -valued random variable on $(\Omega, \mathcal{F}, P)$* if it is a measurable function from $(\Omega, \mathcal{F})$ into $(S, \mathcal{S})$. That is, $X^{-1}(A) \in \mathcal{F}$ for every $A \in \mathcal{S}$.

Special cases:

1. A function $X : \Omega \rightarrow \mathbb{R}$ is called a *real-valued random variable on $(\Omega, \mathcal{F}, P)$* if it is $(\mathbb{R}, \mathcal{B}(\mathbb{R}))$ -valued, where $\mathcal{B}(\mathbb{R})$ is the Borel σ -algebra on $\mathbb{R}$.
2. A function $X : \Omega \rightarrow \mathbb{C}$ is called a *complex-valued random variable on $(\Omega, \mathcal{F}, P)$* if it is $(\mathbb{C}, \mathcal{B}(\mathbb{C}))$ -valued, where $\mathcal{B}(\mathbb{C})$ is the Borel σ -algebra on $\mathbb{C}$.
3. A function $X : \Omega \rightarrow \overline{\mathbb{R}} = \mathbb{R} \cup \{-\infty, \infty\}$ is called a *random variable taking values in the extended real line on $(\Omega, \mathcal{F}, P)$* if it is $(\overline{\mathbb{R}}, \mathcal{B}(\overline{\mathbb{R}}))$ -valued, where $\mathcal{B}(\overline{\mathbb{R}})$ is the Borel σ -algebra on $\overline{\mathbb{R}}$.

A function X as above ($(S, \mathcal{S})$ -valued, real-valued, complex-valued, or extended real-valued) is called simply a *random variable on $(\Omega, \mathcal{F}, P)$* when the codomain and σ -algebra are clear from context.

Definition 2.0.3 (Distribution (law) of a random variable on a probability measure space). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1), and let $X : \Omega \rightarrow S$ be a $(S, \mathcal{S})$ -valued random variable (Definition 2.0.2) where $(S, \mathcal{S})$ is a measurable space.

The *distribution* (or *law*) of X , denoted μ_X , is the measure on $(S, \mathcal{S})$ defined by

$$\mu_X(A) = P(X^{-1}(A)) = P(\{\omega \in \Omega : X(\omega) \in A\})$$

for every $A \in \mathcal{S}$.

In the special case where $S = \mathbb{R}$ and $\mathcal{S} = \mathcal{B}(\mathbb{R})$ is the Borel σ -algebra on $\mathbb{R}$, the distribution of X is uniquely determined by its *cumulative distribution function* $F_X : \mathbb{R} \rightarrow [0, 1]$ given by

$$F_X(t) = \mu_X((-\infty, t]) = P(X \leq t).$$

Definition 2.0.4 (Expectation of a random variable on a probability measure space). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1).

1. If $X : \Omega \rightarrow \mathbb{R}$ is a real-valued random variable, its *expectation* (or *expected value*) is defined by

$$\mathbb{E}[X] := \int_{\Omega} X dP,$$

where the integral on the right-hand side is the Lebesgue integral with respect to the probability measure P .

More explicitly, one defines

- (a) If $X : \Omega \rightarrow [0, \infty)$ is a nonnegative random variable, then

$$\mathbb{E}[X] := \int_{\Omega} X dP,$$

where the integral on the right-hand side is the Lebesgue integral with respect to the probability measure P .

- (b) If $X : \Omega \rightarrow \mathbb{R}$ is a real-valued random variable, then

$$\mathbb{E}[X] := \mathbb{E}[X^+] - \mathbb{E}[X^-],$$

where we decompose $X = X^+ - X^-$ by letting $X^+ = \max(X, 0)$ and $X^- = \max(-X, 0)$; provided that at least one of $\mathbb{E}[X^+]$, $\mathbb{E}[X^-]$ is finite, this is defined as a value in $[-\infty, \infty]$. If both are finite (i.e., X is a P -integrable function), then $\mathbb{E}[X]$ is a well-defined real number.

2. More generally, if $(S, \mathcal{S})$ is a measurable space and if $X : \Omega \rightarrow S$ is an $(S, \mathcal{S})$ -valued random variable (Definition 2.0.2) with distribution (Definition 2.0.3) μ_X on $(S, \mathcal{S})$, and if $g : S \rightarrow \mathbb{R}$ is a measurable function such that $g \in L^1(S, \mathcal{S}, \mu_X)$, then we define the expectation of g :

$$\mathbb{E}[g(X)] = \int_{\Omega} g(X(\omega)) dP(\omega) = \int_S g(x) d\mu_X(x).$$

A random variable X is called *integrable* if $\mathbb{E}[|X|] < \infty$. The space of all integrable real-valued random variables on $(\Omega, \mathcal{F}, P)$ is denoted by $L^1(\Omega, \mathcal{F}, P)$.

Definition 2.0.5 (Cumulative distribution function). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let $X : \Omega \rightarrow (\mathbb{R}, \mathcal{B}(\mathbb{R}))$ be a random variable (Definition 2.0.2), where $\mathcal{B}(\mathbb{R})$ is the Borel σ -algebra on $\mathbb{R}$. Let μ_X be the distribution of X (Definition 2.0.3). The *cumulative distribution function of X* , denoted by F_X , is the function $F_X : \mathbb{R} \rightarrow [0, 1]$ defined by

$$F_X(t) = \mu_X((-\infty, t]) = P(X^{-1}((-\infty, t])).$$

3. INDEPENDENCE (MEASURE-THEORETIC)

Definition 3.0.1 (Independence of σ -algebras). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1).

(♠ TODO: sub-sigma algebra) A finite collection of sub- σ -algebras $\mathcal{G}_1, \mathcal{G}_2, \dots, \mathcal{G}_n$ of $\mathcal{F}$ is called *independent* if for every choice of events $A_i \in \mathcal{G}_i$ ($1 \leq i \leq n$),

$$P\left(\bigcap_{i=1}^n A_i\right) = \prod_{i=1}^n P(A_i).$$

An arbitrary (possibly infinite) collection of sub- σ -algebras $\{\mathcal{G}_j\}_{j \in J}$ of $\mathcal{F}$ is called *independent* if every finite subcollection is independent in the sense above. That is, for any distinct

indices $j_1, j_2, \dots, j_n \in J$ and any choice of events $A_k \in \mathcal{G}_{j_k}$ ($1 \leq k \leq n$),

$$P\left(\bigcap_{k=1}^n A_k\right) = \prod_{k=1}^n P(A_k).$$

A sequence $(\mathcal{G}_n)_{n=1}^\infty$ of sub- σ -algebras is called an *independent sequence* if the set $\{\mathcal{G}_1, \mathcal{G}_2, \dots\}$ is independent.

Definition 3.0.2 (Independence of random variables on a probability measure space). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1).

A finite collection of real-valued random variables $X_1, X_2, \dots, X_n$ on $(\Omega, \mathcal{F}, P)$ is called *independent* if the sub- σ -algebras $\sigma(X_1), \sigma(X_2), \dots, \sigma(X_n)$ are independent (Definition 3.0.1), where for each i ,

$$\sigma(X_i) = \{X_i^{-1}(B) : B \in \mathcal{B}(\mathbb{R})\}$$

is the σ -algebra generated by X_i , and $\mathcal{B}(\mathbb{R})$ is the Borel σ -algebra on $\mathbb{R}$.

An arbitrary (possibly infinite) collection of random variables $\{X_j\}_{j \in J}$ is called *independent* if every finite subcollection is independent.

A sequence $(X_n)_{n=1}^\infty$ of random variables is called an *independent sequence* if the set $\{\sigma(X_1), \sigma(X_2), \dots\}$ consists of independent σ -algebras.

Equivalently, $X_1, X_2, \dots, X_n$ are independent if and only if for every choice of Borel sets $B_1, B_2, \dots, B_n \in \mathcal{B}(\mathbb{R})$,

$$P\left(\bigcap_{i=1}^n \{\omega : X_i(\omega) \in B_i\}\right) = \prod_{i=1}^n P(\{\omega : X_i(\omega) \in B_i\}).$$

4. MOMENTS AND JOINT DISTRIBUTIONS

Definition 4.0.1 (Variance and covariance). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1).

1. Let X be an integrable real-valued random variable with $\mathbb{E}[X]$ finite. The *variance of X* is defined by

$$\text{Var}(X) := \mathbb{E}[(X - \mathbb{E}[X])^2].$$

Equivalently, if $\mathbb{E}[X^2] < \infty$, then

$$\text{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2.$$

2. Let X and Y be integrable real-valued random variables with $\mathbb{E}[XY]$ finite. The *covariance of X and Y* is defined by

$$\text{Cov}(X, Y) := \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])].$$

Equivalently, if $\mathbb{E}[X^2] < \infty$ and $\mathbb{E}[Y^2] < \infty$, then

$$\text{Cov}(X, Y) = \mathbb{E}[XY] - \mathbb{E}[X]\mathbb{E}[Y].$$

The *correlation coefficient of X and Y* (when $\text{Var}(X) > 0$ and $\text{Var}(Y) > 0$) is defined by

$$\rho(X, Y) := \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X)} \sqrt{\text{Var}(Y)}}.$$

A random variable X has zero variance if and only if $X = \mathbb{E}[X]$ almost surely. Covariance is symmetric: $\text{Cov}(X, Y) = \text{Cov}(Y, X)$, and it is bilinear in its arguments.

Definition 4.0.2 (Distribution of a random vector). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let $n \in \mathbb{N}$ and let $X_1, \dots, X_n$ be real-valued random variables (Definition 2.0.2) on $(\Omega, \mathcal{F}, P)$.

The map $\mathbf{X} : \Omega \rightarrow \mathbb{R}^n$ defined by $\mathbf{X}(\omega) = (X_1(\omega), \dots, X_n(\omega))$ is a random variable (Definition 2.0.2) with values in $(\mathbb{R}^n, \mathcal{B}(\mathbb{R}^n))$, where $\mathcal{B}(\mathbb{R}^n)$ is the Borel σ -algebra on $\mathbb{R}^n$. The random variable $\mathbf{X}$ is called a *random vector*.

The *distribution of the random vector $\mathbf{X}$* , denoted $\mu_{\mathbf{X}}$, is the pushforward measure on $(\mathbb{R}^n, \mathcal{B}(\mathbb{R}^n))$ defined by

$$\mu_{\mathbf{X}}(A) = P(\{\omega \in \Omega : (X_1(\omega), \dots, X_n(\omega)) \in A\})$$

for every $A \in \mathcal{B}(\mathbb{R}^n)$. When $n = 1$, this recovers the distribution of a random variable (Definition 2.0.3).

5. INEQUALITIES

Proposition 5.0.1 (Markov's inequality). Let $(X, \mathcal{M}, \mu)$ be a measure space and let $f : X \rightarrow [0, +\infty]$ be a measurable function. For every $t > 0$,

$$\mu(\{x \in X : f(x) \geq t\}) \leq \frac{1}{t} \int_X f d\mu.$$

In particular, if μ is a probability measure (Definition 2.0.1) and $X \geq 0$ is a non-negative random variable, then for every $t > 0$,

$$P(X \geq t) \leq \frac{\mathbb{E}[X]}{t}.$$

Proposition 5.0.2 (Chebyshev's inequality). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1), and let X be a real-valued random variable on $(\Omega, \mathcal{F}, P)$ with finite variance $\text{Var}(X)$. For every $t > 0$,

$$P(|X - \mathbb{E}[X]| \geq t) \leq \frac{\text{Var}(X)}{t^2}.$$

This follows from the (Proposition 5.0.1)Markov inequality applied to the nonnegative random variable $(X - \mathbb{E}[X])^2$.

Proposition 5.0.3 (Cauchy–Schwarz inequality). Let $(X, \mathcal{M}, \mu)$ be a measure space. For any measurable functions $f, g : X \rightarrow \mathbb{C}$ such that $f, g \in L^2(\mu)$,

$$\left| \int_X f \bar{g} d\mu \right| \leq \|f\|_2 \cdot \|g\|_2 = \left(\int_X |f|^2 d\mu \right)^{1/2} \left(\int_X |g|^2 d\mu \right)^{1/2},$$

where $\|\cdot\|_2$ is the L^2 -norm. In particular, if $(\Omega, \mathcal{F}, P)$ is a probability space (Definition 1.0.1) and X, Y are random variables with $\mathbb{E}[X^2] < \infty$ and $\mathbb{E}[Y^2] < \infty$, then

$$|\mathbb{E}[XY]| \leq (\mathbb{E}[X^2])^{1/2}(\mathbb{E}[Y^2])^{1/2},$$

with equality if and only if X and Y are linearly dependent almost surely. This is the special case of Hölder's inequality (Proposition 5.0.6) for conjugate exponents $p = q = 2$.

Proposition 5.0.4 (Lyapunov's inequality). Let $(X, \mathcal{M}, \mu)$ be a measure space with $\mu(X) = 1$ (i.e., μ is a probability measure (Definition 2.0.1)). Let $f : X \rightarrow \mathbb{C}$ be a measurable function. For any $0 < s < t$, if $f \in L^t(\mu)$, then $f \in L^s(\mu)$ and

$$\|f\|_s \leq \|f\|_t,$$

where $\|f\|_r = (\int_X |f|^r d\mu)^{1/r}$ for $r > 0$ is the L^r -norm. The function $r \mapsto \log(\|f\|_r)$ is convex on $(0, \infty)$, and the map $r \mapsto \|f\|_r$ is non-decreasing on $(0, \infty]$.

Proposition 5.0.5 (Minkowski's inequality). Let $(X, \mathcal{M}, \mu)$ be a measure space and let $p \geq 1$. For any measurable functions $f, g : X \rightarrow \mathbb{C}$ such that $f, g \in L^p(\mu)$,

$$\|f + g\|_p \leq \|f\|_p + \|g\|_p,$$

where $\|h\|_p = (\int_X |h|^p d\mu)^{1/p}$ for $h \in L^p(\mu)$ is the L^p -norm. Equality holds if and only if f and g are positively linearly dependent almost everywhere, i.e., there exist constants $a, b \geq 0$, not both zero, such that $af(x) = bg(x)$ for μ -almost every $x \in X$.

Proposition 5.0.6 (Hölder's inequality). Let $(X, \mathcal{A}, \mu)$ be a measure space. Let $p, q \in (1, \infty)$ be conjugate exponents, i.e. $\frac{1}{p} + \frac{1}{q} = 1$. (♠ TODO: conjugate exponents)

Let $f : X \rightarrow \mathbb{R}$ (or $\mathbb{C}$) and $g : X \rightarrow \mathbb{R}$ (or $\mathbb{C}$) be two measurable functions such that f^p and g^q are integrable, i.e.

$$\int_X |f|^p d\mu < \infty \quad \text{and} \quad \int_X |g|^q d\mu < \infty.$$

Then fg is integrable and

$$\int_X |fg| d\mu \leq \left(\int_X |f|^p d\mu \right)^{\frac{1}{p}} \left(\int_X |g|^q d\mu \right)^{\frac{1}{q}}.$$

Proposition 5.0.7 (Jensen's inequality (measure-theoretic)). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let $\varphi : I \rightarrow \mathbb{R}$ be a convex function defined on an interval $I \subseteq \mathbb{R}$. If $X : \Omega \rightarrow I$ is a measurable function such that $X \in L^1(\Omega, \mathcal{F}, P)$ and $\varphi \circ X \in L^1(\Omega, \mathcal{F}, P)$, then

$$\varphi \left(\int_{\Omega} X dP \right) \leq \int_{\Omega} (\varphi \circ X) dP.$$

Proposition 5.0.8 (Chernoff bound). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1), and let $X : \Omega \rightarrow \mathbb{R}$ be a random variable. For every $t \in \mathbb{R}$,

$$P(X \geq t) \leq \inf_{s>0} e^{-st} \mathbb{E}[e^{sX}],$$

provided the expectation on the right-hand side is finite for some $s > 0$.

The bound is obtained by applying the (Proposition 5.0.1)Markov inequality to the non-negative random variable e^{sX} and then optimizing over $s > 0$.

6. MODES OF CONVERGENCE

Definition 6.0.1 (Almost sure convergence). Let $(X, \mathcal{A}, \mu)$ be a measure space. Let T be an ordered commutative monoid, let T^∞ be the ordered monoid with adjoined infinity obtained from T , and let (Y, d_Y) be a metric space valued in T^∞ (typically $Y = \mathbb{R}, \overline{\mathbb{R}},$ or $\mathbb{C}$).

Let $(f_n)_{n \geq 1}$ be a sequence of measurable functions $f_n : X \rightarrow Y$ and let $f : X \rightarrow Y$ be a function. The sequence (f_n) is said to *converge to f almost surely* or to *converge μ -almost everywhere* if

$$\mu\left(\left\{x \in X : d_Y(f_n(x), f(x)) \not\rightarrow 0_T \text{ as } n \rightarrow \infty\right\}\right) = 0.$$

Equivalently, $f_n \rightarrow f$ μ -a.e. if there exists a set $N \in \mathcal{A}$ with $\mu(N) = 0$ such that for all $x \in X \setminus N$, $f_n(x) \rightarrow f(x)$ in the metric d_Y .

Definition 6.0.2 (Convergence in measure). Let $(X, \mathcal{M}, \mu)$ be a measure space. Let T be an ordered commutative monoid, let T^∞ be the ordered monoid with adjoined infinity obtained from T , and let (Y, d) be a metric space valued in T^∞ . Let $\{f_n\}_{n=1}^\infty$ be a sequence of measurable functions $f_n : X \rightarrow Y$, and let $f : X \rightarrow Y$ be a measurable function.

The sequence $\{f_n\}$ is said to *converge in measure (on $(X, \mathcal{M}, \mu)$ into (Y, d)) to f* if for every $\varepsilon \in T_{>0}$,

$$\mu(\{x \in X : d(f_n(x), f(x)) \geq \varepsilon\}) \rightarrow 0 \quad \text{as } n \rightarrow \infty.$$

That is, for every $\varepsilon \in T_{>0}$ and every $\delta > 0$ in $\mathbb{R}$, there exists $N \in \mathbb{N}$ such that for all $n \geq N$,

$$\mu(\{x \in X : d(f_n(x), f(x)) \geq \varepsilon\}) < \delta.$$

Definition 6.0.3. Let $(X, \mathcal{A}, \mu)$ be a measure space and $1 \leq p \leq \infty$.

A sequence of measurable functions $\{f_n\}_{n=1}^\infty$, where each $f_n : X \rightarrow \mathbb{R}$ (or $f_n : X \rightarrow \mathbb{C}$), *converges to a function $f : X \rightarrow \mathbb{R}$ (or $f : X \rightarrow \mathbb{C}$) in the L^p norm* if and only if the sequence $(f_n)_{n=1}^\infty$ converges to f in the sense of (see Definition ??) the extended metric space which (see Remark ??) is the space of all measurable functions equipped with the metric induced by the L^p norm, i.e.

$$\lim_{n \rightarrow \infty} \|f_n - f\|_p = 0.$$

We may thus also say that f is the *limit of $(f_n)_{n=1}^\infty$ for the L^p norm on X* , and write

$$f = \lim_{n \rightarrow \infty} f_n \text{ in } L^p$$

or

$$f_n \rightarrow f \text{ in } L^p.$$

Definition 6.0.4 (Convergence in distribution). (♠ TODO: cumulative distribution) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let $\{X_n\}_{n=1}^\infty$ and X be real-valued random variables (Definition 2.0.2) on $(\Omega, \mathcal{F}, P)$, with cumulative distribution functions F_{X_n} and F_X , respectively. We say that X_n *converges in distribution* to X (denoted $X_n \xrightarrow{d} X$) if for every $t \in \mathbb{R}$ at which F_X is continuous,

$$\lim_{n \rightarrow \infty} F_{X_n}(t) = F_X(t).$$

Proposition 6.0.5 (Relationships between modes of convergence). (♠ TODO: AI generated, verify) Let $(X, \mathcal{A}, \mu)$ be a measure space, let T be an ordered commutative monoid, let T^∞ be the ordered monoid with adjoined infinity obtained from T , and let $(f_n)_{n=1}^\infty$ and f be measurable functions from X to a metric space valued in $T^\infty(S, d)$.

1. If $\mu(X) < \infty$ and $f_n \rightarrow f$ almost surely (i.e., almost everywhere (Definition 6.0.1)), then $f_n \rightarrow f$ almost uniformly.
2. If $f_n \rightarrow f$ almost uniformly, then $f_n \rightarrow f$ in measure.
3. If f_n, f are measurable functions from X to $\mathbb{R}$ and there exists $M > 0$ such that $|f_n(x)| \leq M$ for all n and all $x \in X$, then almost uniform convergence implies convergence in the (Definition 6.0.3) L^p norm for every $1 \leq p < \infty$.
4. If $f_n, f \in L^p(X, \mathcal{A}, \mu)$ for some $1 \leq p < \infty$ and $\|f_n - f\|_p \rightarrow 0$, then there exists a subsequence $(f_{n_k})_{k=1}^\infty$ such that $f_{n_k} \rightarrow f$ almost everywhere (Definition 6.0.1) as $k \rightarrow \infty$.

7. LIMIT THEOREMS

Lemma 7.0.1 (Borel–Cantelli lemmas). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1), and let $\{A_n\}_{n=1}^\infty$ be a sequence of events in $\mathcal{F}$. Define the event

$$\limsup_{n \rightarrow \infty} A_n := \bigcap_{n=1}^\infty \bigcup_{k=n}^\infty A_k = \{\omega \in \Omega : \omega \in A_n \text{ for infinitely many } n\}.$$

First lemma If $\sum_{n=1}^\infty P(A_n) < \infty$, then

$$P(\limsup_{n \rightarrow \infty} A_n) = 0.$$

Second lemma If the events A_n are independent (Definition 1.0.7) and $\sum_{n=1}^\infty P(A_n) = \infty$, then

$$P(\limsup_{n \rightarrow \infty} A_n) = 1.$$

Theorem 7.0.2 (Weak law of large numbers (Khinchine)). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1). Let $\{X_n\}_{n=1}^\infty$ be a sequence of independent and identically distributed (Definition 1.0.8) real-valued random variables on $(\Omega, \mathcal{F}, P)$, with $\mathbb{E}[X_1] = \mu$ and $\text{Var}(X_1) = \sigma^2 < \infty$ (Definition 4.0.1).

Let $S_n = X_1 + \cdots + X_n$ and let $\bar{X}_n = S_n/n$. Then for every $\varepsilon > 0$,

$$\lim_{n \rightarrow \infty} P(|\bar{X}_n - \mu| \geq \varepsilon) = 0.$$

That is, $\bar{X}_n$ converges to μ in probability (Definition 6.0.2). This follows from Chebyshev's inequality (Proposition 5.0.2) applied to $\bar{X}_n$, since $\text{Var}(\bar{X}_n) = \sigma^2/n \rightarrow 0$.

Theorem 7.0.3 (Strong law of large numbers (Kolmogorov)). Let $(\Omega, \mathcal{F}, P)$ be a probability space (Definition 1.0.1), and let $\{X_n\}_{n=1}^\infty$ be a sequence of independent (Definition 3.0.2) and identically distributed real-valued random variables (Definition 2.0.2) on $(\Omega, \mathcal{F}, P)$ such that $\mathbb{E}[|X_1|] < \infty$.

Let $S_n = \sum_{i=1}^n X_i$ for each $n \geq 1$. Then

$$\frac{S_n}{n} \rightarrow \mathbb{E}[X_1] \quad \text{almost surely as } n \rightarrow \infty.$$

8. INTEGRATION THEOREMS (USED IN PROBABILITY)

Theorem 8.0.1 (Monotone convergence theorem). (♠ TODO: AI generated, verify) Let $(X, \mathcal{M}, \mu)$ be a measure space and let $\{f_n\}_{n=1}^\infty$ be a sequence of measurable functions $f_n : X \rightarrow [0, \infty]$ such that $f_1(x) \leq f_2(x) \leq \dots$ for all $x \in X$. Define $f : X \rightarrow [0, \infty]$ by $f(x) = \lim_{n \rightarrow \infty} f_n(x) = \sup_{n \in \mathbb{N}} f_n(x)$. Then each f_n and f are Lebesgue integrable and

$$\lim_{n \rightarrow \infty} \int_X f_n d\mu = \int_X f d\mu.$$

(??)

Theorem 8.0.2 (Dominated convergence theorem). Let $(X, \mathcal{A}, \mu)$ be a measure space.

Let $(f_n)_{n \geq 1}$ be a sequence of measurable functions $f_n : X \rightarrow \overline{\mathbb{R}}$ (or $\mathbb{C}$) and let $f : X \rightarrow \overline{\mathbb{R}}$ (or $\mathbb{C}$) be a function. Suppose that:

1. $f_n \rightarrow f$ almost everywhere (Definition 6.0.1) as $n \rightarrow \infty$, and (♠ TODO: dominated by...)
2. there exists an integrable function $g : X \rightarrow [0, \infty]$ such that $|f_n(x)| \leq g(x)$ for all $n \geq 1$ and all $x \in X$ (i.e. (f_n) is dominated by an integrable function).

Then each f_n (and f) is integrable, and

$$\lim_{n \rightarrow \infty} \int_X f_n d\mu = \int_X f d\mu.$$

Equivalently, $\|f_n - f\|_1 \rightarrow 0$, i.e. $f_n \rightarrow f$ in $L^1(\mu)$.

Theorem 8.0.3 (Bounded convergence theorem). Let $(X, \mathcal{M}, \mu)$ be a finite measure space, i.e., $\mu(X) < \infty$. Let $\{f_n\}_{n=1}^\infty$ be a sequence of measurable functions $f_n : X \rightarrow \mathbb{R}$, and let $f : X \rightarrow \mathbb{R}$ be measurable.

Suppose that:

1. $f_n \rightarrow f$ almost everywhere with respect to μ , and
2. there exists a constant $M < \infty$ such that $|f_n(x)| \leq M$ for all $n \in \mathbb{N}$ and all $x \in X$.

Then each f_n and f are integrable, and

$$\lim_{n \rightarrow \infty} \int_X f_n d\mu = \int_X f d\mu.$$

Theorem 8.0.4 (Fatou's lemma). Let $(X, \mathcal{M}, \mu)$ be a measure space and let $\{f_n\}_{n=1}^\infty$ be a sequence of measurable functions $f_n : X \rightarrow [0, \infty]$. Then

$$\int_X \liminf_{n \rightarrow \infty} f_n d\mu \leq \liminf_{n \rightarrow \infty} \int_X f_n d\mu.$$

(??

9. CONDITIONAL EXPECTATION

Definition 9.0.1 (Conditional expectation). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1), and let $\mathcal{G} \subseteq \mathcal{F}$ be a sub- σ -algebra of $\mathcal{F}$. Let $X : \Omega \rightarrow \mathbb{R}$ be an integrable random variable, i.e., $\mathbb{E}[|X|] < \infty$ (Definition 2.0.4).

A $\mathcal{G}$ -measurable random variable $Y : \Omega \rightarrow \mathbb{R}$ is called the *conditional expectation of X given $\mathcal{G}$* , denoted $\mathbb{E}[X \mid \mathcal{G}]$, if for every $A \in \mathcal{G}$,

$$\int_A Y dP = \int_A X dP.$$

By the Radon–Nikodym theorem, such a Y exists and is unique up to almost sure equality. The measure $\nu(A) = \int_A X dP$ for $A \in \mathcal{G}$ is absolutely continuous with respect to $P|_{\mathcal{G}}$, so $Y = d\nu/d(P|_{\mathcal{G}})$ is the Radon–Nikodym derivative.

Proposition 9.0.2 (Properties of conditional expectation). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1), let $\mathcal{G} \subseteq \mathcal{F}$ be a sub- σ -algebra, and let X be an integrable random variable on $(\Omega, \mathcal{F}, P)$. Then the conditional expectation (Definition 9.0.1) $\mathbb{E}[X \mid \mathcal{G}]$ satisfies the following properties:

1. **(Linearity)** For any $a, b \in \mathbb{R}$ and integrable random variables Y ,

$$\mathbb{E}[aX + bY \mid \mathcal{G}] = a\mathbb{E}[X \mid \mathcal{G}] + b\mathbb{E}[Y \mid \mathcal{G}] \quad \text{almost surely.}$$

2. **(Tower property)** If $\mathcal{H} \subseteq \mathcal{G}$ is a sub- σ -algebra, then

$$\mathbb{E}[\mathbb{E}[X \mid \mathcal{G}] \mid \mathcal{H}] = \mathbb{E}[X \mid \mathcal{H}] \quad \text{almost surely.}$$

3. **(Pulling out what is known)** If Z is $\mathcal{G}$ -measurable and ZX is integrable, then

$$\mathbb{E}[ZX \mid \mathcal{G}] = Z\mathbb{E}[X \mid \mathcal{G}] \quad \text{almost surely.}$$

4. **(Jensen's inequality)** If $\varphi : \mathbb{R} \rightarrow \mathbb{R}$ is a convex function such that $\varphi(X)$ is integrable, then

$$\varphi(\mathbb{E}[X \mid \mathcal{G}]) \leq \mathbb{E}[\varphi(X) \mid \mathcal{G}] \quad \text{almost surely.}$$

10. CHARACTERISTIC FUNCTIONS

Definition 10.0.1 (Characteristic function of a random variable). Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1), and let $X : \Omega \rightarrow \mathbb{R}$ be a real-valued random variable (Definition 2.0.2) with distribution (Definition 2.0.3) μ_X .

The *characteristic function of X* , denoted φ_X , is the function $\varphi_X : \mathbb{R} \rightarrow \mathbb{C}$ defined by

$$\varphi_X(t) = \mathbb{E}[e^{itX}] = \int_{\mathbb{R}} e^{itx} d\mu_X(x)$$

for each $t \in \mathbb{R}$, where $\mathbb{E}$ denotes the expected value (Definition 2.0.4). The integral is well-defined because $|e^{itx}| = 1$ for all $t, x \in \mathbb{R}$.

Theorem 10.0.2 (Characteristic function determines distribution). Let X and Y be two random variables (Definition 2.0.2) defined on possibly different probability measure spaces (Definition 2.0.1), and suppose their distributions (Definition 2.0.3) μ_X and μ_Y have characteristic functions φ_X and φ_Y , respectively.

If $\varphi_X(t) = \varphi_Y(t)$ for all $t \in \mathbb{R}$, then $\mu_X = \mu_Y$.

11. COMMON PROBABILITY DISTRIBUTIONS

11.1. Discrete Distributions.

Definition 11.1.1 (Bernoulli distribution). (♠ TODO: AI generated, verify) Let $p \in [0, 1]$. A real-valued random variable X on a probability measure space (Definition 2.0.1) has the *Bernoulli distribution with parameter p* (denoted $\text{Bern}(p)$) if

$$P(X = 1) = p \quad \text{and} \quad P(X = 0) = 1 - p.$$

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = p$, and the variance (Definition 4.0.1) is $\text{Var}(X) = p(1 - p)$.

Definition 11.1.2 (Binomial distribution). (♠ TODO: AI generated, verify) Let $n \in \mathbb{N}_{\geq 1}$ and $p \in [0, 1]$. A random variable X on a probability measure space (Definition 2.0.1) with codomain $\{0, 1, \dots, n\}$ has the *binomial distribution with parameters n and p* (denoted $\text{Bin}(n, p)$) if for each $k \in \{0, 1, \dots, n\}$,

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}.$$

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = np$, and the variance (Definition 4.0.1) is $\text{Var}(X) = np(1 - p)$.

Definition 11.1.3 (Geometric distribution). (♠ TODO: AI generated, verify) Let $p \in (0, 1]$. A random variable X on a probability measure space (Definition 2.0.1) with codomain $\{1, 2, 3, \dots\}$ has the *geometric distribution with parameter p* (denoted $\text{Geom}(p)$) if for each $k \in \{1, 2, 3, \dots\}$,

$$P(X = k) = (1 - p)^{k-1} p.$$

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = \frac{1}{p}$, and the variance (Definition 4.0.1) is $\text{Var}(X) = \frac{1-p}{p^2}$.

Definition 11.1.4 (Poisson distribution). (♠ TODO: AI generated, verify) Let $\lambda > 0$. A random variable X on a probability measure space (Definition 2.0.1) with codomain $\{0, 1, 2, \dots\}$ has the *Poisson distribution with parameter λ* (denoted $\text{Pois}(\lambda)$) if for each $k \in \{0, 1, 2, \dots\}$,

$$P(X = k) = e^{-\lambda} \frac{\lambda^k}{k!}.$$

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = \lambda$, and the variance (Definition 4.0.1) is $\text{Var}(X) = \lambda$.

11.2. Continuous Distributions.

Definition 11.2.1 (Continuous uniform distribution on an interval). (♠ TODO: AI generated, verify) Let $a, b \in \mathbb{R}$ with $a < b$. A real-valued random variable X on a probability measure space (Definition 2.0.1) has the *continuous uniform distribution on $[a, b]$* (denoted $\text{Unif}(a, b)$) if its distribution (Definition 2.0.3) is the measure on $\mathbb{R}$ given by

$$\mu_X(A) = \int_A \frac{1}{b-a} \mathbf{1}_{[a,b]}(x) dx$$

for every Borel set $A \subseteq \mathbb{R}$.

The probability density function is

$$f(x) = \frac{1}{b-a} \mathbf{1}_{[a,b]}(x).$$

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = \frac{a+b}{2}$, and the variance (Definition 4.0.1) is $\text{Var}(X) = \frac{(b-a)^2}{12}$.

Definition 11.2.2 (Exponential distribution). (♠ TODO: AI generated, verify) Let $\lambda > 0$. A real-valued random variable X on a probability measure space (Definition 2.0.1) has the *exponential distribution with rate λ* (denoted $\text{Exp}(\lambda)$) if its distribution (Definition 2.0.3) is the measure on $\mathbb{R}$ given by

$$\mu_X(A) = \int_A \lambda e^{-\lambda x} \mathbf{1}_{[0,\infty)}(x) dx$$

for every Borel set $A \subseteq \mathbb{R}$.

The probability density function is

$$f(x) = \lambda e^{-\lambda x} \mathbf{1}_{[0,\infty)}(x).$$

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = \frac{1}{\lambda}$, and the variance (Definition 4.0.1) is $\text{Var}(X) = \frac{1}{\lambda^2}$.

Definition 11.2.3 (Normal (Gaussian) distribution on the real line). (♠ TODO: AI generated, verify) Let $\mu \in \mathbb{R}$ and $\sigma > 0$. A real-valued random variable X on a probability measure space (Definition 2.0.1) has the *normal (Gaussian) distribution with mean μ and*

standard deviation σ (denoted $\mathcal{N}(\mu, \sigma^2)$) if its distribution (Definition 2.0.3) is the measure on $\mathbb{R}$ given by

$$\mu_X(A) = \int_A \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx$$

for every Borel set $A \subseteq \mathbb{R}$.

The probability density function is

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}.$$

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = \mu$, and the variance (Definition 4.0.1) is $\text{Var}(X) = \sigma^2$.

Definition 11.2.4 (Gamma distribution). (♠ TODO: AI generated, verify) Let $\alpha > 0$ and $\beta > 0$. A real-valued random variable X on a probability measure space (Definition 2.0.1) has the *gamma distribution with shape α and rate β* (denoted $\text{Gamma}(\alpha, \beta)$) if its distribution (Definition 2.0.3) is the measure on $\mathbb{R}$ given by

$$\mu_X(A) = \int_A \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x} \mathbf{1}_{(0,\infty)}(x) dx$$

for every Borel set $A \subseteq \mathbb{R}$, where $\Gamma(\alpha) = \int_0^\infty t^{\alpha-1} e^{-t} dt$.

The probability density function is

$$f(x) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x} \mathbf{1}_{(0,\infty)}(x).$$

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = \frac{\alpha}{\beta}$, and the variance (Definition 4.0.1) is $\text{Var}(X) = \frac{\alpha}{\beta^2}$.

Definition 11.2.5 (Beta distribution). (♠ TODO: AI generated, verify) Let $\alpha > 0$ and $\beta > 0$. A real-valued random variable X on a probability measure space (Definition 2.0.1) has the *beta distribution with parameters α and β* (denoted $\text{Beta}(\alpha, \beta)$) if its distribution (Definition 2.0.3) is the measure on $\mathbb{R}$ given by

$$\mu_X(A) = \int_A \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1} \mathbf{1}_{(0,1)}(x) dx$$

for every Borel set $A \subseteq \mathbb{R}$, where $B(\alpha, \beta) = \int_0^1 t^{\alpha-1} (1-t)^{\beta-1} dt = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha+\beta)}$.

The probability density function is

$$f(x) = \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1} \mathbf{1}_{(0,1)}(x).$$

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = \frac{\alpha}{\alpha+\beta}$, and the variance (Definition 4.0.1) is $\text{Var}(X) = \frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$.

Definition 11.2.6 (Cauchy distribution). (♠ TODO: AI generated, verify) Let $x_0 \in \mathbb{R}$ and $\gamma > 0$. A real-valued random variable X on a probability measure space (Definition 2.0.1) has the *Cauchy distribution with location parameter x_0 and scale parameter γ* (denoted $\text{Cauchy}(x_0, \gamma)$) if its distribution (Definition 2.0.3) is the measure on $\mathbb{R}$ given by

$$\mu_X(A) = \int_A \frac{1}{\pi\gamma \left[1 + \left(\frac{x-x_0}{\gamma}\right)^2\right]} dx$$

for every Borel set $A \subseteq \mathbb{R}$.

The probability density function is

$$f(x) = \frac{1}{\pi\gamma \left[1 + \left(\frac{x-x_0}{\gamma}\right)^2\right]}.$$

The Cauchy distribution has no finite expectation (Definition 2.0.4) or variance (Definition 4.0.1), since $\int_{-\infty}^{\infty} |x|f(x) dx = \infty$.

11.3. Multivariate Distributions.

Definition 11.3.1 (Multivariate normal (Gaussian) distribution on Euclidean space). (♠ TODO: AI generated, verify) Let $n \in \mathbb{N}_{\geq 1}$, let $\mu \in \mathbb{R}^n$, and let Σ be an $n \times n$ symmetric positive-semidefinite matrix. A random vector $X : \Omega \rightarrow \mathbb{R}^n$ on a probability measure space (Definition 2.0.1) has the *multivariate normal (Gaussian) distribution with mean μ and covariance matrix Σ* (denoted $\mathcal{N}_n(\mu, \Sigma)$) if for every Borel set $A \subseteq \mathbb{R}^n$, its distribution (Definition 4.0.2) is the measure

$$\mu_X(A) = \int_A \frac{1}{(2\pi)^{n/2} \det(\Sigma)^{1/2}} e^{-\frac{1}{2}(x-\mu)^\top \Sigma^{-1}(x-\mu)} dx$$

when Σ is positive-definite. When Σ is singular, the distribution is defined as the weak limit of distributions with nonsingular covariance matrices converging to Σ .

The expectation (Definition 2.0.4) is $\mathbb{E}[X] = \mu$, and the covariance matrix (Definition 4.0.1) is $\text{Cov}(X) = \Sigma$.

12. CENTRAL LIMIT THEOREM

Theorem 12.0.1 (Classical Central Limit Theorem). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let $\{X_n\}_{n=1}^{\infty}$ be a sequence of independent and identically distributed real-valued random variables on $(\Omega, \mathcal{F}, P)$ such that $\mathbb{E}[X_1] = \mu$ and $\text{Var}(X_1) = \sigma^2 \in (0, \infty)$.

Let $S_n = \sum_{i=1}^n X_i$ for each $n \geq 1$. Then the normalized sum

$$\frac{S_n - n\mu}{\sigma\sqrt{n}}$$

converges in distribution (Definition 6.0.4) to a standard normal random variable (Definition 11.2.3) $\mathcal{N}(0, 1)$ as $n \rightarrow \infty$. That is, for every $t \in \mathbb{R}$,

$$\lim_{n \rightarrow \infty} P\left(\frac{S_n - n\mu}{\sigma\sqrt{n}} \leq t\right) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^t e^{-x^2/2} dx.$$

Theorem 12.0.2 (Lindeberg–Feller Central Limit Theorem). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). For each $n \geq 1$, let $\{X_{n,k}\}_{k=1}^{m_n}$ be a triangular array of independent real-valued random variables on $(\Omega, \mathcal{F}, P)$ with $\mathbb{E}[X_{n,k}] = 0$ and $\text{Var}(X_{n,k}) < \infty$.

Let $s_n^2 = \sum_{k=1}^{m_n} \text{Var}(X_{n,k})$ and assume $s_n > 0$ for all n . Suppose the Lindeberg condition holds: for every $\varepsilon > 0$,

$$\lim_{n \rightarrow \infty} \frac{1}{s_n^2} \sum_{k=1}^{m_n} \mathbb{E}[X_{n,k}^2 \mathbf{1}_{\{|X_{n,k}| > \varepsilon s_n\}}] = 0.$$

Then

$$\frac{1}{s_n} \sum_{k=1}^{m_n} X_{n,k} \xrightarrow{d} \mathcal{N}(0, 1) \quad \text{as } n \rightarrow \infty.$$

Theorem 12.0.3 (Lyapunov Central Limit Theorem). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). For each $n \geq 1$, let $\{X_{n,k}\}_{k=1}^{m_n}$ be a triangular array of independent real-valued random variables on $(\Omega, \mathcal{F}, P)$ with $\mathbb{E}[X_{n,k}] = 0$ and $\text{Var}(X_{n,k}) < \infty$.

Let $s_n^2 = \sum_{k=1}^{m_n} \text{Var}(X_{n,k})$ and assume $s_n > 0$ for all n . Suppose there exists $\delta > 0$ such that the Lyapunov condition holds:

$$\lim_{n \rightarrow \infty} \frac{1}{s_n^{2+\delta}} \sum_{k=1}^{m_n} \mathbb{E}[|X_{n,k}|^{2+\delta}] = 0.$$

Then

$$\frac{1}{s_n} \sum_{k=1}^{m_n} X_{n,k} \xrightarrow{d} \mathcal{N}(0, 1) \quad \text{as } n \rightarrow \infty.$$

13. MARTINGALES

Definition 13.0.1 (Filtration of σ -algebras on a probability space). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let I be a linearly ordered index set (typically $I = \mathbb{N}_{\geq 0}$ or $I = [0, \infty)$).

A *filtration of σ -algebras on $(\Omega, \mathcal{F}, P)$ indexed by I* is a family $(\mathcal{F}_t)_{t \in I}$ of sub- σ -algebras of $\mathcal{F}$ such that for all $s, t \in I$ with $s \leq t$,

$$\mathcal{F}_s \subseteq \mathcal{F}_t.$$

The pair $(\Omega, (\mathcal{F}_t)_{t \in I}, P)$ is called a *filtered probability space*. The σ -algebra $\mathcal{F}_t$ represents the information available up to time t .

Definition 13.0.2 (Martingale, submartingale, supermartingale with respect to a filtration). (♠ TODO: AI generated, verify) Let $(\Omega, (\mathcal{F}_t)_{t \in I}, P)$ be a filtered probability space (Definition 13.0.1), where I is a linearly ordered index set. Let $\{X_t\}_{t \in I}$ be a family of real-valued random variables such that for each $t \in I$:

1. X_t is $\mathcal{F}_t$ -measurable (i.e., X_t is *adapted to the filtration* $(\mathcal{F}_t)_{t \in I}$).
2. $\mathbb{E}[|X_t|] < \infty$.

The family $\{X_t\}_{t \in I}$ is called:

1. A *martingale with respect to* $(\mathcal{F}_t)_{t \in I}$ if for all $s \leq t$ in I ,
$$\mathbb{E}[X_t | \mathcal{F}_s] = X_s \quad \text{a.s.}$$
2. A *submartingale with respect to* $(\mathcal{F}_t)_{t \in I}$ if for all $s \leq t$ in I ,
$$\mathbb{E}[X_t | \mathcal{F}_s] \geq X_s \quad \text{a.s.}$$
3. A *supermartingale with respect to* $(\mathcal{F}_t)_{t \in I}$ if for all $s \leq t$ in I ,
$$\mathbb{E}[X_t | \mathcal{F}_s] \leq X_s \quad \text{a.s.}$$

Here $\mathbb{E}[\cdot | \mathcal{F}_s]$ denotes the conditional expectation (Definition 9.0.1).

Theorem 13.0.3 (Optional stopping theorem for martingales). (♠ TODO: AI generated, verify) Let $(\Omega, (\mathcal{F}_n)_{n \geq 0}, P)$ be a filtered probability space (Definition 13.0.1). Let $\{X_n\}_{n \geq 0}$ be a martingale (Definition 13.0.2) with respect to $(\mathcal{F}_n)_{n \geq 0}$.

Let $T : \Omega \rightarrow \mathbb{N}_{\geq 0} \cup \{\infty\}$ be a stopping time with respect to $(\mathcal{F}_n)_{n \geq 0}$, i.e., $\{T = n\} \in \mathcal{F}_n$ for all $n \geq 0$. If one of the following conditions holds:

1. T is bounded almost surely (i.e., there exists $N < \infty$ such that $T \leq N$ a.s.);
2. $\mathbb{E}[T] < \infty$ and $|X_{n+1} - X_n| \leq M$ a.s. for some constant $M > 0$ and all n ;
3. $\mathbb{E}[|X_T|] < \infty$ and $\lim_{N \rightarrow \infty} \mathbb{E}[|X_N| \mathbf{1}_{\{T > N\}}] = 0$,

then the stopped random variable X_T is integrable and

$$\mathbb{E}[X_T] = \mathbb{E}[X_0].$$

Theorem 13.0.4 (Martingale convergence theorem for L^1 -bounded supermartingales). (♠ TODO: AI generated, verify) Let $(\Omega, (\mathcal{F}_n)_{n \geq 0}, P)$ be a filtered probability space (Definition 13.0.1). Let $\{X_n\}_{n \geq 0}$ be a supermartingale (Definition 13.0.2) with respect to $(\mathcal{F}_n)_{n \geq 0}$ such that

$$\sup_{n \geq 0} \mathbb{E}[|X_n|] < \infty.$$

Then there exists an integrable random variable X_∞ such that

$$X_n \rightarrow X_\infty \quad \text{almost surely as } n \rightarrow \infty.$$

If $\{X_n\}_{n \geq 0}$ is a martingale (Definition 13.0.2), then X_∞ satisfies $\mathbb{E}[|X_\infty|] < \infty$, and for each $n \geq 0$,

$$\mathbb{E}[X_\infty | \mathcal{F}_n] = X_n \quad \text{a.s.}$$

when the martingale is uniformly integrable.

14. STOCHASTIC PROCESSES

Definition 14.0.1 (Discrete-time Markov chain on a countable state space). (♠ TODO: AI generated, verify) Let $(\Omega, (\mathcal{F}_n)_{n \geq 0}, P)$ be a filtered probability space (Definition 13.0.1). Let S be a countable set.

A sequence of S -valued random variables $\{X_n\}_{n \geq 0}$ adapted to $(\mathcal{F}_n)_{n \geq 0}$ is called a *discrete-time Markov chain on the state space S* if for all $n \geq 0$ and all states $i, j \in S$,

$$P(X_{n+1} = j \mid X_0, X_1, \dots, X_n = i) = P(X_{n+1} = j \mid X_n = i) \quad \text{a.s.}$$

The function $p_{ij}^{(n)} = P(X_{n+m} = j \mid X_m = i)$ is called the *n -step transition probability from i to j* . When $p_{ij}^{(1)}$ does not depend on m , the chain has *time-homogeneous transitions*, and $P = (p_{ij})_{i,j \in S}$ is called the *transition matrix of the Markov chain*.

Definition 14.0.2 (Poisson process with rate parameter). (♠ TODO: AI generated, verify) Let $\lambda > 0$. A sequence of counting random variables $\{N_t\}_{t \geq 0}$ on a probability measure space (Definition 2.0.1) $(\Omega, \mathcal{F}, P)$ is called a *Poisson process with rate λ* if it satisfies:

1. $N_0 = 0$ almost surely;
2. (Independent increments) For all $0 \leq t_0 < t_1 < \dots < t_n$, the random variables $N_{t_1} - N_{t_0}, N_{t_2} - N_{t_1}, \dots, N_{t_n} - N_{t_{n-1}}$ are independent;
3. (Stationary increments) For all $s < t$, the increment $N_t - N_s$ has a Poisson distribution with parameter $\lambda(t - s)$ (Definition 11.1.4), i.e., for each $k \in \mathbb{N}_{\geq 0}$,

$$P(N_t - N_s = k) = e^{-\lambda(t-s)} \frac{(\lambda(t-s))^k}{k!}.$$

For each $t \geq 0$, $\mathbb{E}[N_t] = \lambda t$ and $\text{Var}(N_t) = \lambda t$.

Definition 14.0.3 (Standard Brownian motion on a probability space). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). A family of real-valued random variables $\{B_t\}_{t \geq 0}$ on $(\Omega, \mathcal{F}, P)$ is called a *standard Brownian motion (Wiener process)* if it satisfies:

1. $B_0 = 0$ almost surely;
2. (Independent increments) For all $0 \leq t_0 < t_1 < \dots < t_n$, the random variables $B_{t_1} - B_{t_0}, B_{t_2} - B_{t_1}, \dots, B_{t_n} - B_{t_{n-1}}$ are independent;
3. (Gaussian increments) For all $s < t$, the increment $B_t - B_s$ has a normal distribution (Definition 11.2.3) $\mathcal{N}(0, t - s)$;
4. (Continuous paths) Almost surely, the map $t \mapsto B_t(\omega)$ is continuous on $[0, \infty)$.

For each $t \geq 0$, $\mathbb{E}[B_t] = 0$ and $\text{Var}(B_t) = t$. The process has stationary increments.

15. ENTROPY AND INFORMATION THEORY

Definition 15.0.1 (Shannon entropy of a discrete random variable). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let $X : \Omega \rightarrow \mathcal{X}$

be a random variable taking values in a finite or countable set $\mathcal{X}$. Let $p_x = P(X = x)$ for each $x \in \mathcal{X}$.

The *Shannon entropy of X* (base e , measured in nats) is defined as

$$H(X) = - \sum_{x \in \mathcal{X}} p_x \log p_x,$$

with the convention $0 \log 0 = 0$. The entropy takes values in $[0, \infty]$. When $\mathcal{X}$ is finite with $|\mathcal{X}| = n$, we have $0 \leq H(X) \leq \log n$, with equality if and only if $p_x = 1/n$ for all x .

Definition 15.0.2 (Differential entropy of a continuous random variable). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let $X : \Omega \rightarrow \mathbb{R}$ be a real-valued random variable with probability density function f with respect to Lebesgue measure on $\mathbb{R}$.

The *differential entropy of X* (base e , measured in nats) is defined as

$$h(X) = - \int_{-\infty}^{\infty} f(x) \log f(x) dx,$$

with the convention $0 \log 0 = 0$. Unlike discrete entropy, differential entropy can be negative and is not invariant under change of variables.

For a random vector $X : \Omega \rightarrow \mathbb{R}^n$ with joint density f , the differential entropy is defined analogously as

$$h(X) = - \int_{\mathbb{R}^n} f(x) \log f(x) dx.$$

Among all random variables on $\mathbb{R}$ with a fixed variance σ^2 , the normal distribution (Definition 11.2.3) $\mathcal{N}(0, \sigma^2)$ maximizes differential entropy, achieving $h(X) = \frac{1}{2} \log(2\pi e \sigma^2)$.

Definition 15.0.3 (Kullback–Leibler divergence between two probability distributions). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let X and Y be random variables on Ω taking values in a finite or countable set $\mathcal{X}$, with probability mass functions $p_x = P(X = x)$ and $q_x = P(Y = x)$ for each $x \in \mathcal{X}$.

The *Kullback–Leibler divergence of Q from P* (also called the *relative entropy of P with respect to Q*), denoted $D_{\text{KL}}(P||Q)$, is defined as

$$D_{\text{KL}}(P||Q) = \sum_{x \in \mathcal{X}} p_x \log \frac{p_x}{q_x},$$

with the conventions $0 \log 0 = 0$, $p_x \log(p_x/0) = +\infty$ for $p_x > 0$, and $D_{\text{KL}}(P||Q) = +\infty$ if there exists x with $p_x > 0$ and $q_x = 0$.

By Gibbs' inequality, $D_{\text{KL}}(P||Q) \geq 0$, with equality if and only if $p_x = q_x$ for all $x \in \mathcal{X}$. The KL divergence is not symmetric.

Theorem 15.0.4 (Data processing inequality for relative entropy). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let X and Y be random variables on Ω taking values in finite or countable sets $\mathcal{X}$ and $\mathcal{Y}$, respectively.

Let P_X and Q_X be two probability distributions on $\mathcal{X}$. Suppose that given $X = x$, the conditional distribution of Y is $W(\cdot | x)$ for each $x \in \mathcal{X}$, independent of whether X follows P_X or Q_X . Let P_Y and Q_Y denote the induced marginal distributions on $\mathcal{Y}$:

$$P_Y(y) = \sum_{x \in \mathcal{X}} P_X(x) W(y | x), \quad Q_Y(y) = \sum_{x \in \mathcal{X}} Q_X(x) W(y | x).$$

Then the Kullback–Leibler divergence (Definition 15.0.3) satisfies

$$D_{\text{KL}}(P_Y \| Q_Y) \leq D_{\text{KL}}(P_X \| Q_X).$$

That is, the KL divergence cannot increase under a stochastic mapping (channel).

Definition 15.0.5 (Mutual information between two random variables). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let X and Y be random variables on Ω taking values in finite or countable sets $\mathcal{X}$ and $\mathcal{Y}$, with joint distribution $p_{XY}(x, y) = P(X = x, Y = y)$ and marginals $p_X(x)$ and $p_Y(y)$.

The *mutual information between X and Y* (base e , measured in nats), denoted $I(X; Y)$, is defined as

$$I(X; Y) = \sum_{x \in \mathcal{X}} \sum_{y \in \mathcal{Y}} p_{XY}(x, y) \log \frac{p_{XY}(x, y)}{p_X(x)p_Y(y)}.$$

Equivalently, $I(X; Y) = D_{\text{KL}}(P_{XY} \| P_X \otimes P_Y)$, where D_{KL} is the Kullback–Leibler divergence (Definition 15.0.3). Thus $I(X; Y) \geq 0$, with equality if and only if X and Y are independent.

For continuous random variables with joint density f_{XY} and marginals f_X, f_Y , the mutual information is defined analogously as

$$I(X; Y) = \int \int f_{XY}(x, y) \log \frac{f_{XY}(x, y)}{f_X(x)f_Y(y)} dx dy.$$

Mutual information satisfies $I(X; Y) = H(X) - H(X | Y) = h(X) - h(X | Y)$ in the discrete and continuous cases respectively.

16. CONDITIONAL PROBABILITY AND BAYES' THEOREM

Definition 16.0.1 (Conditional probability and Bayes' theorem). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let $A, B \in \mathcal{F}$ with $P(B) > 0$.

The *conditional probability of A given B* is defined as

$$P(A | B) = \frac{P(A \cap B)}{P(B)}.$$

If $\{B_1, \dots, B_n\}$ forms a partition of Ω with $P(B_i) > 0$ for each i , then Bayes' theorem states that for any event A with $P(A) > 0$,

$$P(B_j | A) = \frac{P(A | B_j)P(B_j)}{\sum_{i=1}^n P(A | B_i)P(B_i)}.$$

17. MOMENT GENERATING FUNCTIONS

Definition 17.0.1 (Moment generating function of a random variable). (♠ TODO: AI generated, verify) Let $(\Omega, \mathcal{F}, P)$ be a probability measure space (Definition 2.0.1). Let $X : \Omega \rightarrow \mathbb{R}$ be a real-valued random variable.

The *moment generating function of X* , denoted M_X , is the function $M_X : \mathbb{R} \rightarrow [0, \infty]$ defined by

$$M_X(t) = \mathbb{E} [e^{tX}]$$

for each $t \in \mathbb{R}$, whenever the expectation exists.

If $M_X(t)$ is finite for all t in an open neighborhood of 0, then X has moments of all orders and for each $n \geq 1$,

$$\mathbb{E}[X^n] = M_X^{(n)}(0).$$

The *cumulant generating function of X* , denoted K_X , is defined as

$$K_X(t) = \log M_X(t)$$

on the domain where $M_X(t)$ is finite. The n -th derivative $K_X^{(n)}(0)$ gives the n -th cumulant of X .

18. COUPLING AND TOTAL VARIATION DISTANCE

Definition 18.0.1 (Total variation distance between two probability measures). (♠ TODO: AI generated, verify) Let $(S, \mathcal{A})$ be a measurable space. Let μ and ν be two probability measures on $(S, \mathcal{A})$.

The *total variation distance between μ and ν* , denoted $\|\mu - \nu\|_{\text{TV}}$, is defined as

$$\|\mu - \nu\|_{\text{TV}} = \sup_{A \in \mathcal{A}} |\mu(A) - \nu(A)|.$$

When μ and ν are absolutely continuous with respect to a common measure λ , with densities $f = d\mu/d\lambda$ and $g = d\nu/d\lambda$, the total variation distance satisfies

$$\|\mu - \nu\|_{\text{TV}} = \frac{1}{2} \int_S |f(x) - g(x)| d\lambda(x).$$

A *coupling of μ and ν* is a probability measure π on $(S \times S, \mathcal{A} \otimes \mathcal{A})$ whose first marginal is μ and second marginal is ν . The coupling inequality states that for any coupling π ,

$$\|\mu - \nu\|_{\text{TV}} \leq \pi(\{(x, y) : x \neq y\}),$$

and there exists an optimal coupling achieving equality.

19. ERGODIC THEORY (PROBABILITY-THEORETIC VIEW)

Definition 19.0.1 (Measure-preserving transformation on a probability space). (♠ TODO: AI generated, verify) Let $(X, \mathcal{A}, \mu)$ be a probability measure space (Definition 2.0.1). A measurable map $T : X \rightarrow X$ is called a *measure-preserving transformation* if for every $A \in \mathcal{A}$,

$$\mu(T^{-1}(A)) = \mu(A).$$

The system $(X, \mathcal{A}, \mu, T)$ is called a *dynamical system on the probability space*. The transformation T is said to be *ergodic* if for every $A \in \mathcal{A}$ with $T^{-1}(A) = A$, either $\mu(A) = 0$ or $\mu(A) = 1$.

Theorem 19.0.2 (Birkhoff's ergodic theorem for measure-preserving transformations). (♠ TODO: AI generated, verify) Let $(X, \mathcal{A}, \mu)$ be a probability measure space (Definition 2.0.1). Let $T : X \rightarrow X$ be a measure-preserving transformation (Definition 19.0.1). Let $f \in L^1(X, \mathcal{A}, \mu)$.

Define the Birkhoff averages

$$\frac{1}{n} \sum_{k=0}^{n-1} f(T^k x).$$

Then there exists a function $f^* \in L^1(X, \mathcal{A}, \mu)$ such that:

1. The limit $\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{k=0}^{n-1} f(T^k x) = f^*(x)$ exists for μ -almost every $x \in X$.
2. $f^* \circ T = f^*$ almost everywhere.
3. If T is ergodic (Definition 19.0.1), then f^* is constant almost everywhere and

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{k=0}^{n-1} f(T^k x) = \int_X f d\mu \quad \text{for } \mu\text{-a.e. } x.$$